

Advanced Computer Systems Project 2 Report

Ivan Cheng (ECSE)

Rensselaer Polytechnic Institute

ECSE 4320 Advanced Computer Systems

Dr. Tong. Zhang

October 1, 2025

Setup/methodology

The experiments were performed on an Intel Core i7-11850H laptop.

OS: Ubuntu on WSL2

Compiler: g++ with -O3 -std=c++17

SMT enabled, but experiments are single threaded

Experimental Knobs:

Patterns: sequential vs. random

Strides: 64B, 256B, 1024B

R/W ratios: 100R, 100W, 70/30, 50/50

Intensity: 1-8 threads

Page size: 4KB, huge pages are 2MB

Measurement:

Latency from pointer-chasing (ns + cycles per access)

Bandwidth in GB/s

Experiments repeated ≥ 3 times, averages reported

Results & Analysis

Zero-queue baselines

| bytes | pattern | strideB | ns_per_access | cycles_per_access |
|--------|---------|---------|---------------|-------------------|
| 4096 | chase | 64 | 4.10156 | 10.2539 |
| 8192 | chase | 64 | 4.03442 | 10.0861 |
| 16384 | chase | 64 | 4.39453 | 10.9863 |
| 32768 | chase | 64 | 4.15649 | 10.3912 |
| 57344 | chase | 64 | 5.73556 | 14.3389 |
| 110592 | chase | 64 | 5.7147 | 14.2867 |
| 212992 | chase | 64 | 5.68824 | 14.2206 |
| 409600 | chase | 64 | 5.83875 | 14.5969 |

| | | | | |
|--------------|-------|----|---------|---------|
| 786432 | chase | 64 | 5.68345 | 14.2086 |
| 1511424 | chase | 64 | 6.14649 | 15.3662 |
| 2912256 | chase | 64 | 5.99178 | 14.9795 |
| 5619712 | chase | 64 | 5.9945 | 14.9862 |
| 1082982 4 | chase | 64 | 6.11753 | 15.2938 |
| 2088550 4 | chase | 64 | 6.9635 | 17.4087 |
| 4027187 2 | chase | 64 | 8.54962 | 21.3741 |
| 7765606 4 | chase | 64 | 8.77782 | 21.9446 |
| 1.5E+08 | chase | 64 | 8.9482 | 22.3705 |
| 2.89E+0 8 | chase | 64 | 9.10721 | 22.768 |
| 5.57E+0 8 | chase | 64 | 8.96056 | 22.4014 |
| 1.07E+0 9 | chase | 64 | 8.95298 | 22.3825 |

Based on the data, the range is 4ns-9ns per access.

This is expected as the smaller working sets that fit in L1 are ~4ns.

Larger sets that spill into L2/L3 are ~5-7 ns.

DRAM are ~8-9ns.

This matches hierarchical behavior as latency in ns/access increases as working set grows.

Pattern & granularity sweep

| bytes | pattern | strideB | GBps | ns_per_ele m |
|-------|---------|---------|--------------|-----------------|
| 65536 | random | 64 | 0.77121 3 | 15.5599 |

| | | | | |
|----------|--------|----|---------|----------|
| 196608 | random | 64 | 3.00216 | 3.99712 |
| 569344 | random | 64 | 10.9888 | 1.09202 |
| 1667072 | random | 64 | 12.5344 | 0.957367 |
| 4894720 | random | 64 | 11.7649 | 1.01998 |
| 14385152 | random | 64 | 14.2327 | 0.843129 |
| 42278912 | random | 64 | 13.0147 | 0.922033 |
| 1.24E+08 | random | 64 | 13.087 | 0.916943 |
| 3.65E+08 | random | 64 | 13.2944 | 0.902633 |
| 1.07E+09 | random | 64 | 12.3263 | 0.97353 |

| bytes | pattern | strideB | GBps | ns_per_element |
|----------|---------|---------|----------|----------------|
| 65536 | random | 256 | 0.637303 | 18.8293 |
| 196608 | random | 256 | 2.30041 | 5.21647 |
| 569344 | random | 256 | 7.21805 | 1.6625 |
| 1667072 | random | 256 | 13.6199 | 0.881066 |
| 4894720 | random | 256 | 20.7756 | 0.577602 |
| 14385152 | random | 256 | 22.1875 | 0.540845 |
| 42278912 | random | 256 | 23.8921 | 0.502258 |
| 1.24E+08 | random | 256 | 23.9641 | 0.500748 |

| | | | | |
|----------|--------|-----|---------|----------|
| 3.65E+08 | random | 256 | 23.5472 | 0.509615 |
| 1.07E+09 | random | 256 | 20.484 | 0.585824 |

| bytes | pattern | strideB | GBps | ns_per_element |
|----------|---------|---------|----------|----------------|
| 65536 | random | 1024 | 0.799761 | 15.0045 |
| 196608 | random | 1024 | 2.18588 | 5.48977 |
| 569344 | random | 1024 | 5.70994 | 2.1016 |
| 1667072 | random | 1024 | 16.7789 | 0.715186 |
| 4894720 | random | 1024 | 40.9448 | 0.293078 |
| 14385152 | random | 1024 | 60.4955 | 0.198362 |
| 42278912 | random | 1024 | 79.8904 | 0.150206 |
| 1.24E+08 | random | 1024 | 89.7967 | 0.133635 |
| 3.65E+08 | random | 1024 | 83.411 | 0.143866 |
| 1.07E+09 | random | 1024 | 75.6237 | 0.15868 |

| bytes | pattern | strideB | GBps | ns_per_element |
|-------|---------|---------|----------|----------------|
| 65536 | seq | 64 | 0.645746 | 18.5832 |

| | | | | |
|----------|-----|----|---------|----------|
| 196608 | seq | 64 | 2.2521 | 5.32837 |
| 569344 | seq | 64 | 5.23615 | 2.29176 |
| 1667072 | seq | 64 | 11.6452 | 1.03047 |
| 4894720 | seq | 64 | 14.1643 | 0.847199 |
| 14385152 | seq | 64 | 13.0277 | 0.921116 |
| 42278912 | seq | 64 | 12.9246 | 0.928463 |
| 1.24E+08 | seq | 64 | 13.222 | 0.907575 |
| 3.65E+08 | seq | 64 | 13.3334 | 0.899993 |
| 1.07E+09 | seq | 64 | 12.1813 | 0.985116 |

| bytes | pattern | strideB | GBps | ns_per_element |
|----------|---------|---------|----------|----------------|
| 65536 | seq | 256 | 0.606565 | 19.7835 |
| 196608 | seq | 256 | 2.31183 | 5.1907 |
| 569344 | seq | 256 | 7.62628 | 1.57351 |
| 1667072 | seq | 256 | 15.5769 | 0.770373 |
| 4894720 | seq | 256 | 20.119 | 0.596452 |
| 14385152 | seq | 256 | 22.2857 | 0.538463 |
| 42278912 | seq | 256 | 23.1877 | 0.517516 |
| 1.24E+08 | seq | 256 | 23.9606 | 0.500822 |

| | | | | |
|----------|-----|-----|---------|----------|
| 3.65E+08 | seq | 256 | 23.7591 | 0.50507 |
| 1.07E+09 | seq | 256 | 20.7368 | 0.578681 |

| bytes | pattern | strideB | GBps | ns_per_element |
|----------|---------|---------|----------|----------------|
| 65536 | seq | 1024 | 0.759495 | 15.8 |
| 196608 | seq | 1024 | 2.69244 | 4.45692 |
| 569344 | seq | 1024 | 6.57441 | 1.82526 |
| 1667072 | seq | 1024 | 16.0398 | 0.748138 |
| 4894720 | seq | 1024 | 36.6554 | 0.327373 |
| 14385152 | seq | 1024 | 57.8052 | 0.207594 |
| 42278912 | seq | 1024 | 79.7598 | 0.150452 |
| 1.24E+08 | seq | 1024 | 83.0983 | 0.144407 |
| 3.65E+08 | seq | 1024 | 84.4556 | 0.142087 |
| 1.07E+09 | seq | 1024 | 59.402 | 0.202013 |

In my data, we see that in the 64B dataset the bandwidth rises from 0.85GB/s to 13GB/s at large sizes with ns/element falling close to 0.9ns. This case shows the best efficiency baseline. At the 256B dataset the bandwidth goes up to 24GB/s at peak and latency goes further down to 0.5 ns/element. This shows that there are fewer memory requests per byte of allocated space. At 1024B dataset bandwidth peaks at 90GB/s with ns/element at 0.13ns. Demand is massively reduced and GBps is inflated as a lot of data is skipped. Large strides reduce cache pressure and memory seems to be fast.

Between the random and sequential access datasets we can see that the sequential access at small sizes GBps is modest but increase quickly. As working set DRAM bandwidth is fully utilized, and ns/elements drops steadily until bandwidth saturates.

In random access dataset GBps is significantly lower than sequential at the same sizes. In addition, ns/elements is higher which shows stalls. Growth is slower as size increases. This matches expectations as sequential curves climb higher and flatten out while random curves plateau at a lower ceiling.

Read/Write mix sweep

| bytes | rw | strideB | GBps | ns_per_element |
|----------|--------|---------|---------|----------------|
| 5.37E+08 | 50R50W | 64 | 1.99261 | 20.0742 |
| 5.37E+08 | 70R30W | 64 | 1.49436 | 26.7674 |
| 5.37E+08 | 100R | 64 | 1.06547 | 37.5421 |
| 5.37E+08 | 100W | 64 | 3.08045 | 12.9851 |

In this experiment, the read/write ratio was varied across 100%R, 100%W, 70R/30W, 50R/50W while keeping the stride and size constant. The results show how the system responds to read and write dominated workloads.

In 100%R, throughput was the highest as reads benefit from hardware prefetching and can be pipelined efficiently. Latency per element was correspondingly the lowest, indicating that the memory hierarchy is optimized for streaming read access.

In 100%W performance was lowest. Writes require cache line invalidations and store buffer handling and can't be prefetched in the same way as reads. As a result, throughput dropped significantly and per element latency increased.

In mixed ratios, writes degraded performance compared to only reads. Overall, this is all expected as the results align with hardware expectations as it is expected that writes degrade performance.

Intensity sweep

| bytes | strideB | threads | GBps | ns_per_element |
|-------|---------|---------|------|----------------|
|-------|---------|---------|------|----------------|

| | | | | |
|----------|---------|---|---------|----------|
| 1.07E+09 | 64 (T1) | 1 | 13.319 | 0.900969 |
| 1.07E+09 | 64 (T2) | 1 | 13.2801 | 0.903605 |
| 1.07E+09 | 64 (T2) | 2 | 23.3406 | 0.514126 |
| 1.07E+09 | 64 (T4) | 1 | 13.4487 | 0.892278 |
| 1.07E+09 | 64 (T4) | 2 | 23.3375 | 0.514193 |
| 1.07E+09 | 64 (T4) | 4 | 28.279 | 0.424343 |
| 1.07E+09 | 64 (T8) | 1 | 13.3369 | 0.899757 |
| 1.07E+09 | 64 (T8) | 2 | 23.9071 | 0.501944 |
| 1.07E+09 | 64 (T8) | 4 | 28.6569 | 0.418747 |
| 1.07E+09 | 64 (T8) | 8 | 24.3062 | 0.493701 |

Based on my results, GB/s raised from $T = 1 \rightarrow T = 4$

As $T = 8$ it started to flatten out and there was small improvement
ns/element decreased at first and went back up at $T = 8$.

This is expected as the memory system can keep up to 4 threads of requests efficiently.

At $T = 8$, DRAM queues are saturated and more concurrency increases queuing,

Working-set size sweep

| | | | |
|-------|-------------|-------------------|-----------------------|
| bytes | stride B | ns_per_acce ss | cycles_per_acce ss |
|-------|-------------|-------------------|-----------------------|

| | | | |
|--------------|----|---------|---------|
| 4096 | 64 | 4.08936 | 10.2234 |
| 8192 | 64 | 4.03442 | 10.0861 |
| 16384 | 64 | 4.02527 | 10.0632 |
| 32768 | 64 | 4.02374 | 10.0594 |
| 57344 | 64 | 5.6536 | 14.134 |
| 110592 | 64 | 5.67898 | 14.1975 |
| 212992 | 64 | 5.67674 | 14.1918 |
| 409600 | 64 | 5.66882 | 14.1721 |
| 786432 | 64 | 5.66241 | 14.156 |
| 1511424 | 64 | 6.368 | 15.92 |
| 2912256 | 64 | 5.97188 | 14.9297 |
| 5619712 | 64 | 6.01906 | 15.0477 |
| 1082982 4 | 64 | 6.05957 | 15.1489 |
| 2088550 4 | 64 | 6.96222 | 17.4055 |
| 4027187 2 | 64 | 8.54217 | 21.3554 |
| 7765606 4 | 64 | 8.80131 | 22.0033 |
| 1.5E+08 | 64 | 8.8509 | 22.1272 |
| 2.89E+0 8 | 64 | 8.95922 | 22.398 |
| 5.57E+0 8 | 64 | 8.94094 | 22.3524 |
| 1.07E+0 9 | 64 | 8.92155 | 22.3039 |

Based on my data, my data properly shows the cache locality transitions.

In the L1 cache region, my data is around 4KB to 32KB with ~4ns per access.

This is the expected L1 cache latency.

In the L2 cache region, my data is around 57KB to 512KB with latency rising to 5.6ns per access. This is the L2 hit latency and is larger than L1 but faster than DRAM.

In the L3 cache region, latency jumps to ~6.3ns. From 2MB to 20MB latency is around 15-17 cycles.

In the DRAM region, which is 40 MB and beyond, latency rises to ~8.5-9ns (21-22 cycles).

My results matches how expected cycling and ns per access jumps discretely when going from different levels of cache memory.

Cache-miss impact

| bytes | pattern | strideB | elapsed_ s |
|----------|---------|---------|---------------|
| 65536 | seq | 64 | 2.80E-06 |
| 159744 | seq | 64 | 7.07E-06 |
| 385024 | seq | 64 | 3.45E-05 |
| 925696 | seq | 64 | 0.00012 |
| 2236416 | seq | 64 | 0.00032 |
| 5398528 | seq | 64 | 0.000915 |
| 13041664 | seq | 64 | 0.002804 |
| 31506432 | seq | 64 | 0.006806 |
| 76124160 | seq | 64 | 0.017435 |
| 1.84E+08 | seq | 64 | 0.040932 |
| 4.44E+08 | seq | 64 | 0.099455 |

| | | | |
|----------|-----|----|----------|
| 1.07E+09 | seq | 64 | 0.236811 |
|----------|-----|----|----------|

| bytes | pattern | strideB | elapsed_s |
|----------|---------|---------|-----------|
| 65536 | seq | 4096 | 3.67E-07 |
| 159744 | seq | 4096 | 4.67E-07 |
| 385024 | seq | 4096 | 1.37E-06 |
| 925696 | seq | 4096 | 4.30E-06 |
| 2236416 | seq | 4096 | 1.40E-05 |
| 5398528 | seq | 4096 | 5.69E-05 |
| 13041664 | seq | 4096 | 0.000147 |
| 31506432 | seq | 4096 | 0.00035 |
| 76124160 | seq | 4096 | 0.000793 |
| 1.84E+08 | seq | 4096 | 0.001962 |
| 4.44E+08 | seq | 4096 | 0.005369 |
| 1.07E+09 | seq | 4096 | 0.013014 |

| bytes | pattern | strideB | elapsed_s |
|--------|---------|---------|-----------|
| 65536 | random | 64 | 3.20E-06 |
| 159744 | random | 64 | 8.43E-06 |

| | | | |
|----------|--------|----|----------|
| 385024 | random | 64 | 4.93E-05 |
| 925696 | random | 64 | 0.000163 |
| 2236416 | random | 64 | 0.000519 |
| 5398528 | random | 64 | 0.002121 |
| 13041664 | random | 64 | 0.007572 |
| 31506432 | random | 64 | 0.023785 |
| 76124160 | random | 64 | 0.073893 |
| 1.84E+08 | random | 64 | 0.227151 |
| 4.44E+08 | random | 64 | 0.625386 |
| 1.07E+09 | random | 64 | 1.6695 |

| bytes | pattern | strideB | elapsed_s |
|----------|---------|---------|-----------|
| 65536 | random | 4096 | 4.00E-07 |
| 159744 | random | 4096 | 4.00E-07 |
| 385024 | random | 4096 | 9.33E-07 |
| 925696 | random | 4096 | 2.57E-06 |
| 2236416 | random | 4096 | 1.64E-05 |
| 5398528 | random | 4096 | 5.18E-05 |
| 13041664 | random | 4096 | 0.000158 |

| | | | |
|--------------|--------|------|----------|
| 3150643 2 | random | 4096 | 0.000431 |
| 7612416 0 | random | 4096 | 0.001296 |
| 1.84E+0 8 | random | 4096 | 0.004 |
| 4.44E+0 8 | random | 4096 | 0.011638 |
| 1.07E+0 9 | random | 4096 | 0.032187 |

Based on my data, latency is lowest in my sequential data access in 64B as as the hardware prefetcher brings in data smoothly. Miss rate is minimized until cache capacity is exceeded. In my sequential 4KB data, entire cache lines are skipped between accesses. This results in more cache misses as the prefetcher only works in between ~1-2KB regions.

In my random cache data accesses the prefetcher can't predict random jumps so caches for the random 64B is worse compared to the sequential 64B but still better than the 4KB stride. The random 4KB accesses are random and sparse. This inflates the runtime as there are more misses.

TLB-miss impact

| bytes | huge | page_stride B | random | elapsed_ s | huge_grante d |
|--------------|------|------------------|--------|---------------|------------------|
| 4194304 | 0 | 64 | 1 | 0.00198 9 | 0 |
| 4194304 | 1 | 2097152 | 1 | 1.87E-06 | 1 |
| 1022771 2 | 0 | 64 | 1 | 0.00522 6 | 0 |
| 1022771 2 | 1 | 2097152 | 1 | 5.37E-06 | 1 |
| 2493235 2 | 0 | 64 | 1 | 0.01566 1 | 0 |

| | | | | | |
|--------------|---|---------|---|--------------|---|
| 2493235 2 | 1 | 2097152 | 1 | 2.13E-05 | 1 |
| 6078464 0 | 0 | 64 | 1 | 0.04204 8 | 0 |
| 6078464 0 | 1 | 2097152 | 1 | 4.97E-05 | 1 |
| 1.48E+0 8 | 0 | 64 | 1 | 0.117143 | 0 |
| 1.48E+0 8 | 1 | 2097152 | 1 | 0.000118 | 1 |
| 3.61E+0 8 | 0 | 64 | 1 | 0.34556 6 | 0 |
| 3.61E+0 8 | 1 | 2097152 | 1 | 0.00031 | 1 |
| 8.81E+0 8 | 0 | 64 | 1 | 0.97092 3 | 0 |
| 8.81E+0 8 | 1 | 2097152 | 1 | 0.00073 | 1 |
| 2.15E+0 9 | 0 | 64 | 1 | 2.59477 | 0 |
| 2.15E+0 9 | 1 | 2097152 | 1 | 0.00200 2 | 1 |

From my data experiment, we can see that normal pages show superlinear growth over .002s at 4MB to ~2.6s at 2GB. This is what is expected as TLB misses start to dominate as the dataset grows. Huge pages remain stable, staying in the micro/millisecond range at 2GB. The speedup grows with dataset size. This matches theory as with a TLB reach with 4KB pages access past around ~6MB start risking misses.

Latency Percentiles

| Workload | QD | P50(ms) | p95(ms) | p99(ms) | p99.9(ms) |
|----------|----|---------|---------|---------|-----------|
|----------|----|---------|---------|---------|-----------|

| | | | | | |
|--------------|----|------|------|------|------|
| 4k rand read | 8 | .029 | .051 | .057 | .108 |
| 4k rand read | 64 | .030 | .052 | .060 | .135 |

The table shows the percentile latencies for a 4KiB random read workload at QD = 8 and QD = 64. The difference in the latencies although small shows the effect of queuing while throughput scales with QD, tail latency grows. For SLA, systems often set performance guarantees at p99/p99.9. This means that if average latency remains below the SLA threshold, tail violations can cause requests to miss the deadline.

Data was collected through DiskSpd with powershell lines

```
.\diskspd.exe -c1G -b4K -r -d30 -t1 -o8 -L -w0 -Rtext -Sh testfile.dat > results_qd8.txt
```

```
.\diskspd.exe -c1G -b4K -r -d30 -t1 -o64 -L -w0 -Rtext -Sh testfile.dat > results_qd64.txt
```

Conclusion

This project successfully shows cache and memory performance across the hierarchy of a modern CPU. Each experiment produces results that generally show results that are consistent with theory. Despite a few anomalies that occurred such as small working set sizes sometimes showing unstable bandwidth and latency numbers. The results provide a general clear idea on how caching, bandwidth, and address translations impact the performance of a CPU.

Reference

[1] OpenAI, "ChatGPT," ChatGPT, 2025. <https://chatgpt.com/>

[2]microsoft, "Releases · microsoft/diskspd," GitHub, Jun. 13, 2024.
<https://github.com/microsoft/diskspd/releases>