

Advanced Computer Systems Project 3 Report

Ivan Cheng (ECSE)

Rensselaer Polytechnic Institute

ECSE 4320 Advanced Computer Systems

Dr. Tong. Zhang

October 1, 2025

Setup/methodology

Test Environment:

The experiments were performed on an Intel Core i7-11850H laptop. The laptop runs on Tiger Lake-H, 8 cores/16 threads, 2.5 Ghz base.

Memory: 32 GB Ram DDR4

Storage: Samsung MZVL2512HCJQ-00BL7 - 512GB Pcie TLC M.2 2280 SSD

OS: Windows

Benchmark tool: Microsoft DiskSpd

A 16 GiB file (testfile.dat) was created to ensure workloads were large enough to pass windows file caching and engage the SSD.

Benchmark:

1. Zero-queue baselines
Measure QD=1 latency for: (a) 4 KiB random reads & writes, (b) 128 KiB sequential reads & writes.
2. Block-size sweep (pattern fixed)
Sweep 4 KiB→256 KiB random and sequential runs.
Produce IOPS/MB/s and average latency
3. Read/Write mix
100%R, 100%W, 70/30, 50/50 ratios
4. Queue-depth sweep
QD values: 1, 2, 4, 8, 16, 32, 64
Single throughput vs. latency trade-off curve. Identify the knee & relation to Little's Law
5. Tail Latency
DiskSpd's percentile reporting (-Rxml) was enabled to capture p50/p95/p99/p99.9 latencies.

PSL commands:

-Sh disables OS cache, -L enables latency statistics

```
.\diskspd.exe -c1G -d10 -b4k -r -t1 -o1 -w100 -Sh -L -Rtext C:\Temp\4k_rand_write.txt  
.\diskspd.exe -c1G -d10 -b128k -s -t1 -o1 -w0 -Sh -L -Rtext C:\Temp\128k_seq_read.txt  
.\diskspd.exe -c1G -d10 -b128k -s -t1 -o1 -w100 -Sh -L -Rtext C:\Temp\128k_seq_write.txt  
.\diskspd.exe -c1G -d10 -b4k -r -t1 -o4 -w0 -Sh -L -Rtext C:\Temp\rand_4k.txt  
.\diskspd.exe -c1G -d10 -b16k -r -t1 -o4 -w0 -Sh -L -Rtext C:\Temp\rand_16k.txt  
.\diskspd.exe -c1G -d10 -b32k -r -t1 -o4 -w0 -Sh -L -Rtext C:\Temp\rand_32k.txt  
.\diskspd.exe -c1G -d10 -b64k -r -t1 -o4 -w0 -Sh -L -Rtext C:\Temp\rand_64k.txt  
.\diskspd.exe -c1G -d10 -b128k -r -t1 -o4 -w0 -Sh -L -Rtext C:\Temp\rand_128k.txt  
.\diskspd.exe -c1G -d10 -b256k -r -t1 -o4 -w0 -Sh -L -Rtext C:\Temp\rand_256k.txt  
.\diskspd.exe -c1G -d10 -b4k -s -t1 -o4 -w0 -Sh -L -Rtext C:\Temp\seq_4k.txt
```

```

.\diskspd.exe -c1G -d10 -b16k -s -t1 -o4 -w0 -Sh -L -Rtext C:\Temp\seq_16k.txt
.\diskspd.exe -c1G -d10 -b32k -s -t1 -o4 -w0 -Sh -L -Rtext C:\Temp\seq_32k.txt
.\diskspd.exe -c1G -d10 -b64k -s -t1 -o4 -w0 -Sh -L -Rtext C:\Temp\seq_64k.txt
.\diskspd.exe -c1G -d10 -b128k -s -t1 -o4 -w0 -Sh -L -Rtext C:\Temp\seq_128k.txt
.\diskspd.exe -c1G -d10 -b256k -s -t1 -o4 -w0 -Sh -L -Rtext C:\Temp\seq_256k.txt
.\diskspd.exe -c1G -d10 -b4k -r -t1 -o4 -w0 -Sh -L -Rtext C:\Temp\mix_4k_100R.txt
.\diskspd.exe -c1G -d10 -b4k -r -t1 -o4 -w100 -Sh -L -Rtext C:\Temp\mix_4k_100W.txt
.\diskspd.exe -c1G -d10 -b4k -r -t1 -o4 -w30 -Sh -L -Rtext C:\Temp\mix_4k_70R30W.txt
.\diskspd.exe -c1G -d10 -b4k -r -t1 -o4 -w50 -Sh -L -Rtext C:\Temp\mix_4k_50R50W.txt
.\diskspd.exe -c1G -d10 -b4k -r -t1 -o1 -w0 -Sh -L -Rtext C:\Temp\qd1_read.txt
.\diskspd.exe -c1G -d10 -b4k -r -t1 -o2 -w0 -Sh -L -Rtext C:\Temp\qd2_read.txt
.\diskspd.exe -c1G -d10 -b4k -r -t1 -o4 -w0 -Sh -L -Rtext C:\Temp\qd4_read.txt
.\diskspd.exe -c1G -d10 -b4k -r -t1 -o8 -w0 -Sh -L -Rtext C:\Temp\qd8_read.txt
.\diskspd.exe -c1G -d10 -b4k -r -t1 -o16 -w0 -Sh -L -Rtext C:\Temp\qd16_read.txt

```

Zero-queue Baselines

Accurate average & percentile latencies for 4 KiB random and 128 KiB sequential (R/W)

Test Case	IOPS (k)	Bandwidth (MB/s)	Avg Latency (ms)	P50 Latency (ms)	p95 Latency (ms)	p99 Latency (ms)	P99.9 Latency (ms)
4 KiB random read, QD=1	27,595	107.8	.036	.031	.038	.051	.139
4 KiB random write, QD=1	27,545	107.6	.036	.032	.039	.053	0.117
128 KiB sequential read, QD=1	14,416	1802	.069	.059	.071	.096	0.283
128 KiB sequential write, QD=1	11,960	1495	.083	.062	.080	.118	.644

Based on my data in my zero-queue baseline experiments, data is expected as when comparing read vs writes reads are faster at larger blocks. In the random 4KiB experiments, data shows low latency and limited bandwidth. At higher blocks (128KiB) experiment, there is higher bandwidth with higher latency. At QD = 1 we can see there is no queuing delay.

$p_{50} \approx$ pure device service time

Tail ($p_{99/99.9}$) represents infrequent device/OS events.

In terms of SLA implications, $p_{99.9} \leq .1\text{ms}$ will miss as I measured .139ms in my 4 KiB write.

$P_{99} \leq .06\text{ms}$ for 4KiB reads will most likely be fine as I measured .051ms.

For 128KiB reads my $p_{99.9}$ was .283ms so reads < 200ms would miss.

For the 128KiB write I noticed that the $p_{99.9}$ jumped to .644 this spike may have been due to background OS scheduling, side delays or other factors.

Block-size Sweep

Block Size	Access Pattern	IOPS	Bandwidth (MB/s)	Avg Latency (ms)
4KiB	Random	66,554	260	.036
4KiB	Seq	65527	256	.027
16KiB	Random	65231	1019	.041
16KiB	Seq	64133	1002	.041
32KiB	Random	66329	2073	.035
32KiB	Seq	66310	2072	.035
64KiB	Random	61144	3822	.049
64KiB	Seq	43993	2750	.085
128KiB	Random	49904	6238	.077
128KiB	Seq	36540	4567	.105
256KiB	Random	24667	6167	.161
256KiB	Seq	22250	5563	.174

Based on my data, numbers match expected storage scaling behavior. Small blocks have high IOPS with low throughput MB/s and low latency. When compared to medium blocks they have lower IOPS with higher throughput in MB/s. In the larger block cases IOPS drops even further and MB/s saturates as the SSD hits the maximum throughput it can deliver. Latency increases as operations require more transfer of data. When comparing random and sequential access

patterns, the sequential data usually hits higher throughput at higher block sizes since data is written in read/write order and the ssd can use internal prefetching and coalescing. In smaller blocks random and sequential processes perform similarly.

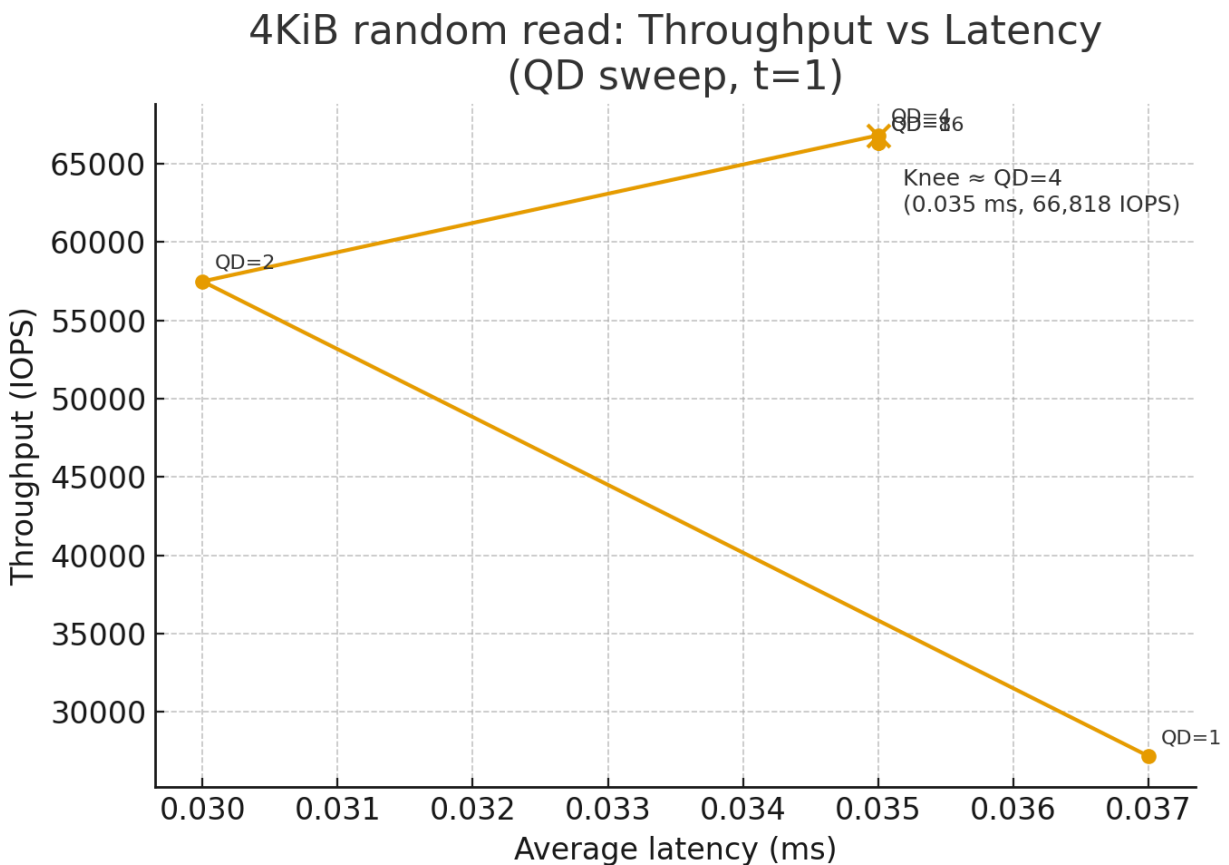
Read/write mix

Mix Ratio	IOPS (k)	Bandwidth (MB/s)	Avg Latency (ms)	P95 Latency (ms)	P99 Latency (ms)
100% Read	66613	260.2	.036	.051	.057
100% Write	59592	232.8	.032	.040	.052
70R/30W	50027	195.4	.068	.098	.967
50R/50W	63127	246.6	.046	.089	.126

Based on the data received through running my experiments, Pure read/write cases deliver the highest throughput and lowest latencies. In the mixed workload cases both runs have reduced IOPS and higher tail latencies because writes contend with reads and flash must erase in addition to reading. My 70/30 case is worse than my 50/50 case in terms of tail latency most likely due to controller scheduling prioritizing writes differently.

Queue-depth sweep

QD	IOPS (k)	Bandwidth (MB/s)	Avg Latency (ms)	P95 Latency (ms)	P99 Latency (ms)
QD = 1	27173	106.14	.037	.037	.047
QD = 2	57490	224.57	.030	.040	.055
QD = 4	66818	261.01	.035	.051	.055
QD = 8	66313	259.04	.035	.051	.056
QD = 16	66327	259.09	.035	.051	.056



With my data, it is observed that going from QD 1 to QD 2 over doubles the IOPS. At QD 4 it goes up to 66k IOPS and afterwards starts to saturate and QD 8 and 16 show no improvement. QD 1 has an average latency of .037 ms and improves to .03 at QD = 2. At QD ≥ 4 average latency stabilizes to .035ms with P95 and P99 latency around .05ms & .056ms respectively. The SSD saturates around QD = 4 which means additional parallelism does not yield more throughput.

In relation to Little's Law $\text{Throughput} \approx \text{Concurrency}/\text{Latency}$ it is predicted that when QD goes past 4 throughput would increase. However measured results stay around 66k. This divergence indicates the device limit for this workload has been reached and additional I/O sits queued. Effective service rate caps out at QD = 4 and throughput no longer follows Little Law.

Anomalies/Limitations

During the course of this project a few anomalies and limitations were encountered:

1. WSL environment issues.

Initially, the experiments were conducted under WSL. However there were issues with getting results with installing FIO onto WSL and issues with how WSL handles system calls and disks I/O. As a result testing was moved to DiskSpd on native Windows.

2. PSL write execution results

When attempting to automate multiple tests at once in one file using PSL scripts an anomaly occurred where write results would output 0 IOPS, 0 MB/s essentially all outputs of write were showing 0. To resolve this each test command was ran manually in Powershell instead of relying on the automated loop.

Conclusion

This project evaluated storage performance under a range of benchmarks using Microsoft's DiskSpd. The results highlighted differences between random vs. sequential access, block-size scaling, read/write mixes, and queue depth behavior. Overall, while DiskSpd provided reproducible measurements, there were anomalies throughout this project such as automation in PSL and environment setup in WSL.

References

- [1] OpenAI, "ChatGPT," ChatGPT, 2025. <https://chatgpt.com/>
- [2] microsoft, "Releases · microsoft/diskspd," GitHub, Jun. 13, 2024. <https://github.com/microsoft/diskspd/releases>
- [3] "Intel® Core™ i7-11850H Processor (24M Cache, up to 4.80 GHz)," Intel, 2025. <https://www.intel.com/content/www/us/en/products/sku/213799/intel-core-i711850h-processor-24m-cache-up-to-4-80-ghz/specifications.html?wapkw=i7-11850H>